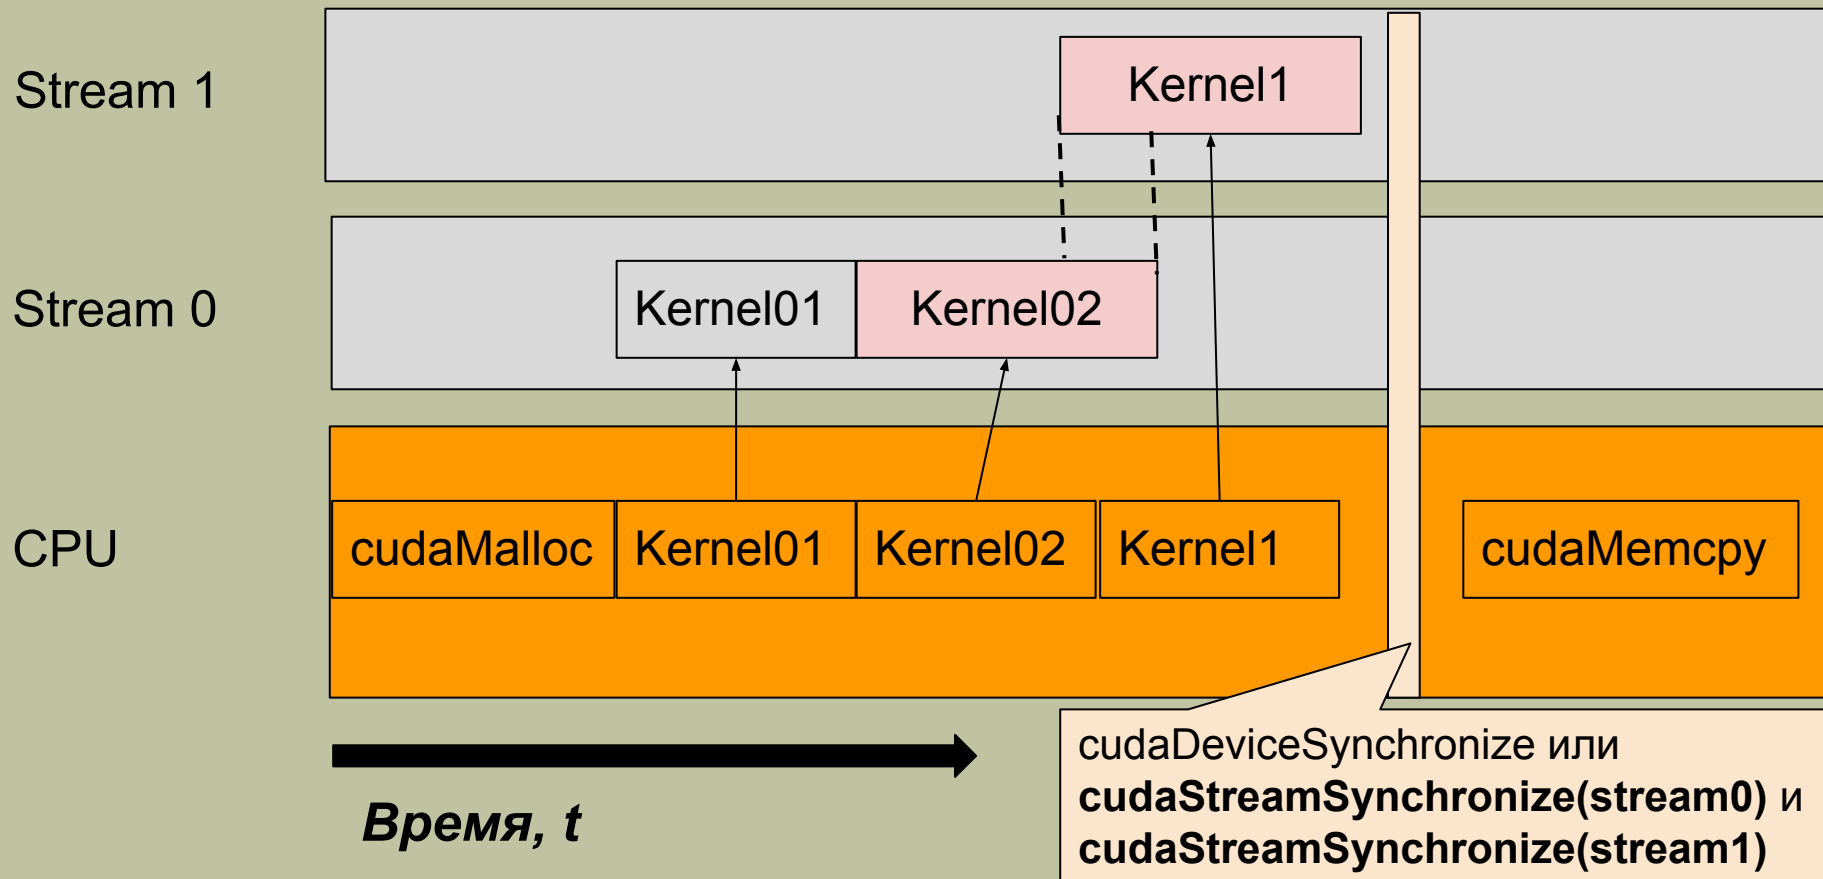


Лекция 9

Параллелизм по задачам.

- Потоки *CUDA (CUDA Stream)*.
- Одновременное выполнение ядер.
- Одновременное копирование и выполнение ядра.
- Использование нескольких GPU.

Потоки CUDA (*CUDA Streams*)



```
#include <stdio.h>
#include <malloc.h>
```

```
__global__ void gTest(int step){
    int n=threadIdx.x + blockIdx.x*blockDim.x;
    printf("kernel, threadIdx: %d\t%d\n", step, n);
}

int main(){
    int NS=3;
    for (int i = 0; i < NS; i++)
        gTest<<< i+1, 32>>>(i);
    cudaDeviceSynchronize();

    return 0;
}
```

Timeline View



1x



3 warnings, 15 messages

0s

+226ms

+226.05ms

+226.1ms

+226.15ms

+226.2ms

+226.25ms

[0] (17.4%)

[0] (17.4%)

2 processes hidden...



151, 8305[93.390 ms]

CPU (6)

Processes (1)

[8305] lab9b

CUDA HW (NVIDIA GeForce RTX 20

100.0% Kernels

100.0% gTest

100.0% gTest(int)

Threads (8)

gTest

gTest

gTest

gTest

gTest

gTest

gTest

gTest

gTest

Events View

Name



#	Name	Start	Duration	GPU	Context
1	gTest	0.226009s	51.840 μ s	GPU 0	Stream 13
2	gTest	0.226062s	66.463 μ s	GPU 0	Stream 13
3	gTest	0.226129s	110.176 μ s	GPU 0	Stream 13

Description:

```
__global__ void gTest(int step){  
    int n=threadIdx.x + blockIdx.x*blockDim.x;  
    printf("kernel, thIdx: %d\t%d\n", step, n);  
}
```

```
int main(){  
    int NS = 3;  
    cudaStream_t *streams;  
  
    streams = (cudaStream_t*)calloc(NS, sizeof(cudaStream_t));  
  
    for (int i = 0; i < NS; i++)  
        cudaStreamCreate(&streams[i]);
```

```
for (int i = 0; i < NS; i++)  
    gTest<<< i+1, 32, 0, streams[i] >>>(i);
```

```
cudaDeviceSynchronize();
```

```
for (int i = 0; i < num_streams; i++)  
    cudaStreamDestroy(streams[i]);
```

```
free(streams);
```

```
return 0;
```

```
}
```

Project 4 x Report 1 x Report 1 x Project 6 x Report 1 x Report 2 x Report 3 x Report 4 x Report 5 x Report 6 x

Timeline View



1x



3 warnings, 15 messages



Events View

Name



#	Name	Start	Duration	GPU	Context
1	gTest	0.181247s	242.399 μ s	GPU 0	Stream 13
2	gTest	0.181254s	235.294 μ s	GPU 0	Stream 14
3	gTest	0.181257s	232.222 μ s	GPU 0	Stream 15

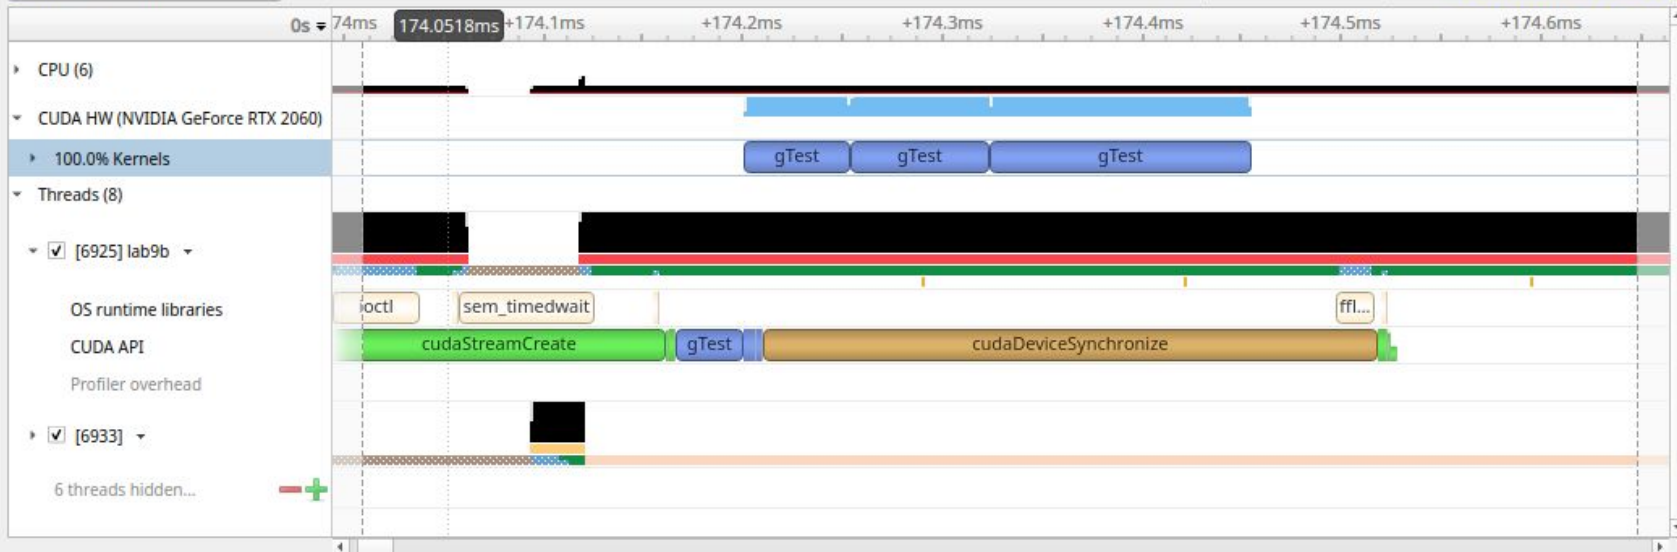
Description:

Project 4 × Report 1 × Report 1 × Project 6 × Report 1 × Report 2 × Report 3 × Report 4 × Report 5 × Report 6 × Report 7 ×

Timeline View

1x

4 warnings, 14 messages



Events View

Name

#	Name	Start	Duration	GPU	Context
1	gTest	0.174202s	51.775 µs	GPU 0	Stream 13
2	gTest	0.174254s	69.696 µs	GPU 0	Stream 13
3	gTest	0.174325s	129.791 µs	GPU 0	Stream 13

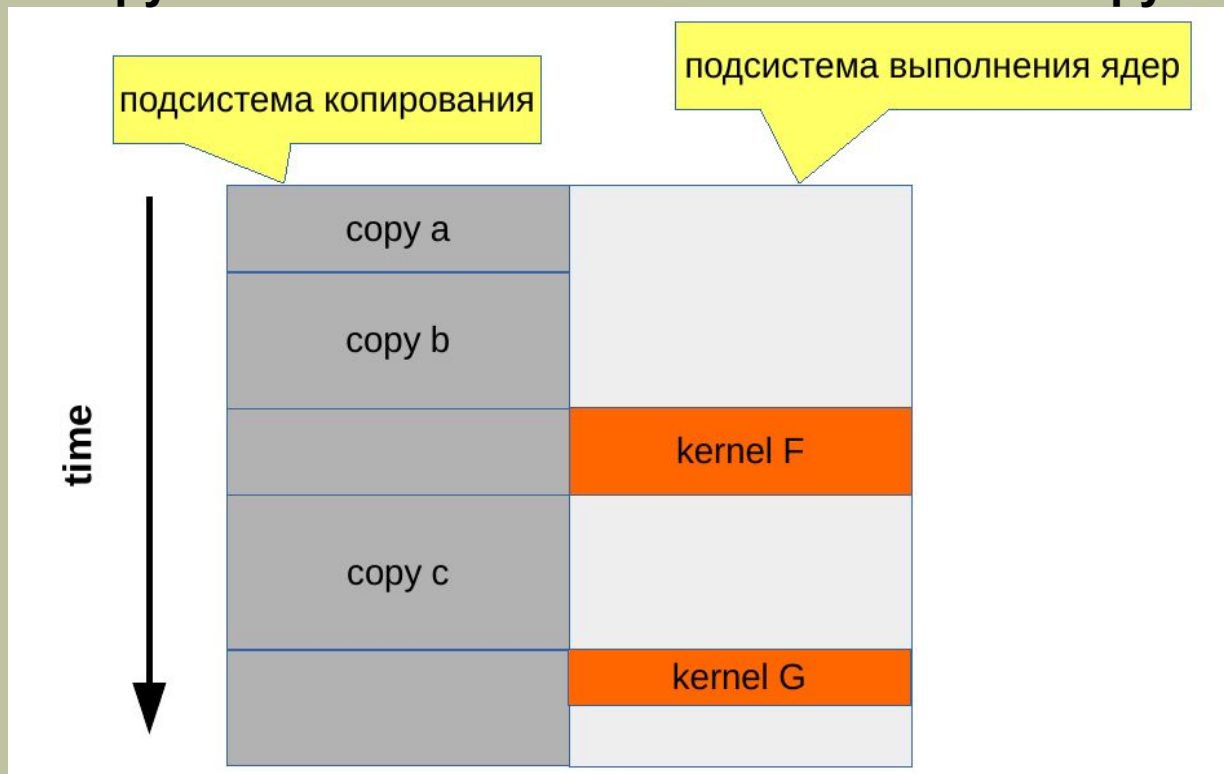
Description:

Device 0: "NVIDIA GeForce RTX 2060"

CUDA Driver Version / Runtime Version 12.0 / 11.1

CUDA Capability Major/Minor version number: 7.5

.....
Concurrent copy and kernel execution: Yes with 3 copy engine(s)



Потоки CUDA и разрешение зависимостей при распараллеливании копирования и выполнения

Очередь копир.

Очередь выполн.

stream0, copy a	
stream0, copy b	
блокировка	kernel0
stream0, copy c	
stream1, copy a	
stream1, copy b	
блокировка	kernel1
stream1, copy c	

Очередь копир.

Очередь выполн.

stream0, copy a	
stream0, copy b	
stream1, copy a	kernel0
stream1, copy b	
stream0, copy c	kernel1
stream1, copy c	

```
#define N (1024*1024)
#define FULL_DATA_SIZE (N*20)

__global__ void kernel(int* a, int* b, int* c){
    int idx=threadIdx.x+blockIdx.x*blockDim.x;
    if(idx<N){
        int idx1=(idx+1)%256;
        int idx2=(idx+2)%256;
        float as=(a[idx]+a[idx1]+a[idx2])/3.0f;
        float bs=(b[idx]+b[idx1]+b[idx2])/3.0f;
        c[idx]=(as+bs)/2;
    }
}
```

```
int main(){
    cudaDeviceProp prop;
    int whichDevice;

    cudaGetDevice(&whichDevice);

    cudaGetDeviceProperties(&prop, whichDevice);
    if(!prop.deviceOverlap){
        printf("Device does not support overlapping\n");
        return 0;
    }
}
```

```
int *host_a, *host_b, *host_c;  
int *dev_a, *dev_b, *dev_c;
```

```
cudaMalloc( (void**)&dev_a, N*sizeof(int)) ;  
cudaMalloc( (void**)&dev_b, N*sizeof(int)) ;  
cudaMalloc( (void**)&dev_c, N*sizeof(int)) ;
```

**Выделение
прикрепленной памяти
(pinned memory) на
хосте.**

```
cudaMallocHost( (void**)&host_a, FULL_DATA_SIZE*sizeof(int)) ;  
cudaMallocHost( (void**)&host_b, FULL_DATA_SIZE*sizeof(int)) ;  
cudaMallocHost( (void**)&host_c, FULL_DATA_SIZE*sizeof(int)) ;
```

```
for(int i=0; i<FULL_DATA_SIZE;i++){  
    host_a[i]=rand();  
    host_b[i]=rand();  
}
```

```
cudaStream_t stream;
```

```
cudaStreamCreate(&stream);
```

```
for(int i=0; i<FULL_DATA_SIZE; i+=N){
```

```
    cudaMemcpyAsync(dev_a, host_a+i, N*sizeof(int),  
                    cudaMemcpyHostToDevice, stream);
```

```
    cudaMemcpyAsync(dev_b, host_b+i, N*sizeof(int),  
                    cudaMemcpyHostToDevice, stream);
```

```
    kernel<<<N/256, 256, 0, stream>>>(dev_a, dev_b, dev_c);
```

```
    cudaMemcpyAsync(host_c+i, dev_c, N*sizeof(int),  
                    cudaMemcpyDeviceToHost, stream);
```

```
}
```

```
cudaStreamSynchronize( stream );
```

```
cudaFreeHost(host_a);  
cudaFreeHost(host_b);  
cudaFreeHost(host_c);
```

```
cudaFree(dev_a);  
cudaFree(dev_b);  
cudaFree(dev_c);
```

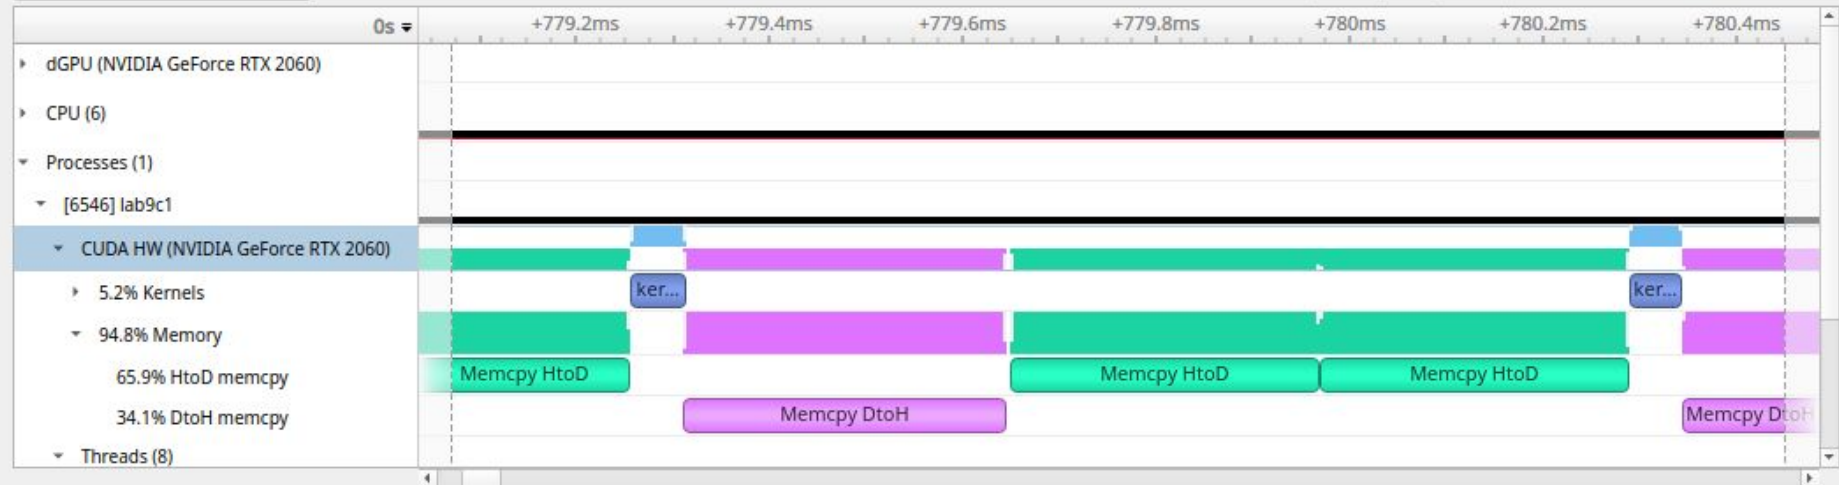
```
cudaStreamDestroy(stream);
```

```
return 0;  
}
```

Timeline View

1x

3 warnings, 15 messages



Events View

Name ▾



#	Name	Start	Duration	GPU	Context	Description:
1	Memcpy HtoD	0.778937s	317.341 μ s	GPU 0	Stream 13	
2	kernel	0.779259s	52.831 μ s	GPU 0	Stream 13	
3	Memcpy DtoH	0.779314s	329.341 μ s	GPU 0	Stream 13	
4	Memcpy HtoD	0.779652s	316.797 μ s	GPU 0	Stream 13	
5	Memcpy HtoD	0.779971s	316.861 μ s	GPU 0	Stream 13	
6	kernel	0.780292s	52.704 μ s	GPU 0	Stream 13	
7	Memcpy DtoH	0.780346s	328.861 μ s	GPU 0	Stream 13	


```
cudaStream_t stream0, stream1;  
    cudaStreamCreate(&stream0);  
    cudaStreamCreate(&stream1);
```

```
for(int i=0; i<FULL_DATA_SIZE;...
```

```
    cudaStreamSynchronize( stream0) ;  
    cudaStreamSynchronize( stream1);
```


```
.....  
    cudaStreamDestroy(stream0);  
    cudaStreamDestroy(stream1);  
.....
```

```
for(int i=0; i<FULL_DATA_SIZE; i+=N*2){  
    cudaMemcpyAsync(dev_a0, host_a+i, N*sizeof(int),  
                    cudaMemcpyHostToDevice, stream0);  
    cudaMemcpyAsync(dev_a1, host_a+i+N, N*sizeof(int),  
                    cudaMemcpyHostToDevice, stream1);  
    cudaMemcpyAsync(dev_b0, host_b+i, N*sizeof(int),  
                    cudaMemcpyHostToDevice, stream0);  
    cudaMemcpyAsync(dev_b1, host_b+i+N, N*sizeof(int),  
                    cudaMemcpyHostToDevice, stream1);  
    kernel<<<N/256, 256, 0, stream0>>>(dev_a0, dev_b0, dev_c0);  
    kernel<<<N/256, 256, 0, stream1>>>(dev_a1, dev_b1, dev_c1);  
    cudaMemcpyAsync(host_c+i, dev_c0, N*sizeof(int),  
                    cudaMemcpyDeviceToHost, stream0);  
    cudaMemcpyAsync(host_c+i+N, dev_c1, N*sizeof(int),  
                    cudaMemcpyDeviceToHost, stream1);  
}
```

Timeline View

1x

3 warnings, 15 messages



Events View

Name

#	Name	Start	Duration	GPU	Context	Description:
1	Memcpy HtoD	0.923951s	418.972 μ s	GPU 0	Stream 13	Begins: 0.923951s
2	Memcpy HtoD	0.924371s	324.798 μ s	GPU 0	Stream 13	Ends: 0.92437s (+418.972 μ s)
3	kernel	0.924706s	56.479 μ s	GPU 0	Stream 13	HtoD memcopy 4,194,304 bytes
4	Memcpy DtoH	0.924769s	464.636 μ s	GPU 0	Stream 13	Source memory kind: Pinned
5	Memcpy HtoD	0.925491s	449.852 μ s	GPU 0	Stream 13	Destination memory kind: Device
						Throughput: 10.0109 GiB/s
						Correlation ID: 120
						Stream: Stream 13

Использование нескольких GPU

```
#include <stdio.h>
#define REAL float

__global__ void initFun(int* nf, int devnum){
    int n=threadIdx.x + blockIdx.x*blockDim.x;
    nf[n]*=10;
}
```

```
int main(int argc, char* argv){  
    if(argc<4){  
        fprintf(stderr, "USAGE: <prog_name> <size_of_array> <num_of_devices>"  
            "<device_indices>\n");  
        return -1;  
    }  
    int N=atoi(argv[1]);
```

```
int* info_devs=(int*)calloc(argc-2, sizeof(int));
info_devs[0]=atoi(argv[2]);
for(int i=1;i<argc-2;i++){
    info_devs[i]=atoi(argv[i+2]);
}
fprintf(stderr,"num of devices: %d\n",info_devs[0]);
for(int i=1;i<argc-2;i++)
    fprintf(stderr,"i_d=%d\n",info_devs[i]);

int** nfd=(int**)calloc(info_devs[0], sizeof(int*));
int** nfh=(int**)calloc(info_devs[0], sizeof(int*));
```

```
cudaStream_t* streams;  
streams=(cudaStream_t*)calloc(info_devs[0], sizeof(cudaStream_t));
```

```
for(int i=0;i<info_devs[0];i++){  
    cudaSetDevice(info_devs[i+1]);  
    cudaStreamCreate(&streams[i]);
```

```
    cudaMalloc((void**)&nfd[i], (N/info_devs[0])*sizeof(int));  
    cudaMallocHost((void**)&nfh[i], (N/info_devs[0])*sizeof(int));
```

```
for(int n=0;n<N/info_devs[0]; n++)  
    nfh[i][n]=n+i*N/info_devs[0];
```

```
cudaMemcpyAsync(nfd[i],nfh[i],  
                (N/info_devs[0])*sizeof(int),  
                cudaMemcpyHostToDevice, streams[i]);  
  
initFun<<<N/info_devs[0]/32, 32, 0, streams[i]>>>(nfd[i],i);  
  
cudaMemcpyAsync(nfh[i],nfd[i],  
                (N/info_devs[0])*sizeof(int),  
                cudaMemcpyDeviceToHost,streams[i]);  
}
```



```
for(int i=0;i<info_devs[0];i++){
    cudaSetDevice(info_devs[i+1]);
    cudaStreamSynchronize(streams[i]);

    for(int n=0;n<N/info_devs[0];n++)
        fprintf(stdout,"nfh[%d][%d]=%d\n",i,n, nfh[i][n]);

    cudaFree(nfd[i]);
    cudaFreeHost(nfh[i]);
    cudaStreamDestroy(streams[i]);
    cudaDeviceReset();
}

return 0;
}
```

```
> ./main6 1024 2 0 1 > tmp61
```

```
num of devices: 2
```

```
i_d=0
```

```
i_d=1
```

```
> vim tmp61
```

```
nfh[0][0]=0
```

```
nfh[0][1]=10
```

```
nfh[0][2]=20
```

```
nfh[0][3]=30
```

```
.....
```

```
nfh[0][510]=5100
```

```
nfh[0][511]=5110
```

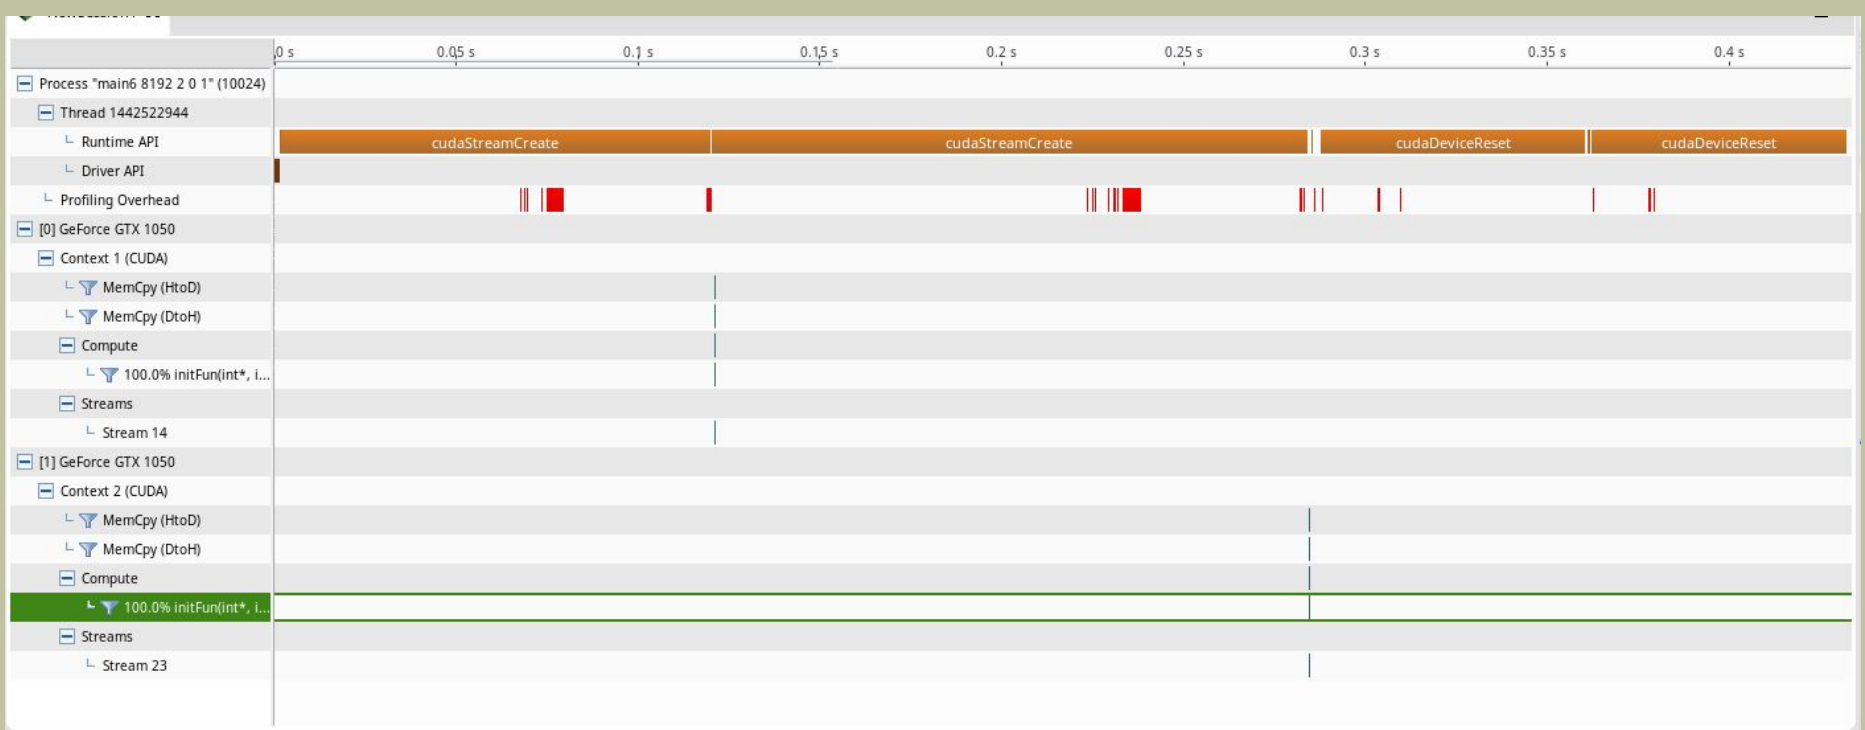
```
nfh[1][0]=5120
```

```
nfh[1][1]=5130
```

```
.....
```

```
nfh[1][510]=10220
```

```
nfh[1][511]=10230
```



Analysis GPU Details (Summary) CPU Details OpenACC Details Console Settings

Name	Invocations	Avg. Duration	Regs	Static SMem	Avg. Dynamic SMem
initFun(int*, int)	1	2.24 μ s	8	0	0

- [-] Process "main6 8192 2 0 1" (11886)
 - [-] Thread 2036029248
 - └ Runtime API
 - └ Driver API
 - └ Profiling Overhead
 - [-] [0] GeForce GTX 1050
 - [-] Context 1 (CUDA)
 - └ MemCpy (HtoD)
 - └ MemCpy (DtoH)
 - [-] Compute
 - └ 100.0% initFun(int*, i..
 - [-] Streams
 - └ Stream 14
 - [-] [1] GeForce GTX 1050
 - [-] Context 2 (CUDA)**
 - └ MemCpy (HtoD)
 - └ MemCpy (DtoH)
 - [-] Compute
 - └ 100.0% initFun(int*, i..
 - [-] Streams
 - └ Stream 23

